



---

# Classification:

## More Classification Techniques


**CS 418. Introduction to Data Science**

© 2018 by Gonzalo A. Bello




# $k$ -Nearest Neighbors Classifier Preliminaries


- A  $k$ -nearest neighbors classifier assigns class labels to observations based on the class labels of the  $k$  “**most similar**” observations ( $k$  nearest neighbors).
- “If it walks like a duck, quacks like a duck, and looks like a duck, then it’s probably a duck.”
- *Example:*



1-Nearest Neighbor



2-Nearest Neighbors



3-Nearest Neighbors



# $k$ -Nearest Neighbors Classifier Algorithm

- **Algorithm for  $k$ -nearest neighbors classifier:**
  1. Choose a **proximity metric** and a **number of nearest neighbors  $k$** .
  2. For each observation  $z$  in the test set:
    1. Compute the **distance** or **similarity** between  $z$  and every observation in the training set.
    2. Select the  $k$  observations **most similar** to  $z$  ( **$k$  nearest neighbors**).
    3. Assign a class label to  $z$  based on the class labels of its  **$k$  nearest neighbors**.
      - Use **majority voting** or **weighted majority voting**.
      - If there is a **tie** between classes, randomly choose one of them.
- **How to choose  $k$ ?**
  - If  $k$  is **too small**, then the classifier may be susceptible to **overfitting** due to noise.
  - If  $k$  is **too large**, then the classifier may misclassify  $z$  by considering observations that are **not similar** to  $z$ .



# **$k$ -Nearest Neighbors Classifier**

## **Distance Metrics**

- **Distance metrics** measure the **dissimilarity** between two observations.
  - For **continuous** attributes:

- **Euclidean distance:**

$$d(x, y) = \sqrt{\sum_{i=1}^p (x_i - y_i)^2}$$

where  $p$  is the number of attributes and  $x_i$  and  $y_i$  are the  $i$ th attributes of observations  $x$  and  $y$ , respectively.

- **Manhattan distance:**

$$d(x, y) = \sum_{i=1}^p |x_i - y_i|$$

- The **Euclidean** and the **Manhattan** distances can be generalized by the **Minkowski distance**:

$$d(x, y) = \left( \sum_{i=1}^p |x_i - y_i|^r \right)^{1/r}$$

where  $r$  is a parameter.



# ***k*-Nearest Neighbors Classifier**

## **Similarity Metrics**

- **Similarity metrics** measure the **similarity** between two observations.
  - For **binary** attributes:

- **Jaccard coefficient:**

$$S(x, y) = \frac{N_{11}}{N_{01} + N_{10} + N_{11}}$$

- **Simple matching coefficient:**

$$S(x, y) = \frac{N_{00} + N_{11}}{N_{00} + N_{01} + N_{10} + N_{11}}$$

where  $N_{11}$  is the number of attributes where both  $x$  and  $y$  have a value of 1,  $N_{00}$  is the number of attributes where both  $x$  and  $y$  have a value of 0,  $N_{01}$  is the number of attributes where  $x$  has a value of 0 and  $y$  has a value of 1, and  $N_{10}$  is the number of attributes where  $x$  has a value of 1 and  $y$  has a value of 0.

- For more **distance** and **similarity metrics** available on Python, see [sklearn.neighbors.DistanceMetric](#).



# $k$ -Nearest Neighbors Classifier Advantages and Disadvantages

- What are some of the **advantages** of  $k$ -nearest neighbors?
  - Does not require training a model.
  - **Nonparametric**. Makes no assumptions about the probability distribution of the data.
  - Applicable to **categorical and numeric attributes**.
  - **Flexibility**. Can produce decision boundaries of any shape.
- What are some of the **disadvantages** of  $k$ -nearest neighbors?
  - Classifying each observation from the test set can be **computationally expensive**.
  - **Easy to overfit** due to noise or high dimensionality.
  - **Sensitive to redundant or irrelevant attributes**.
  - **Sensitive to the scale** of the data.
  - **No descriptive value**. Does not provide a model of the relationship between the class label and the attributes.



# Naïve Bayes Classifier Preliminaries (I)

- A **Naïve Bayes classifier** is a **probabilistic model** that uses **Bayes' Theorem** for classification.

The **probability** of observing class label  **$y$**  given attribute value  **$X$**  is

$$P(y|X) = \frac{P(X|y) P(y)}{P(X)} = \frac{P(X|y) P(y)}{P(X|y) P(y) + P(X|\neg y) P(\neg y)}$$

- Example:** Suppose that we want to classify emails as **spam** or **non-spam** based on the appearance of the word “free” in the email. Using **Bayes' Theorem** we can compute:

$$P(\text{spam}|\text{free}) = \frac{P(\text{free}|\text{spam}) P(\text{spam})}{P(\text{free}|\text{spam}) P(\text{spam}) + P(\text{free}|\neg\text{spam}) P(\neg\text{spam})}$$

$$P(\neg\text{spam}|\text{free}) = \frac{P(\text{free}|\neg\text{spam}) P(\neg\text{spam})}{P(\text{free}|\text{spam}) P(\text{spam}) + P(\text{free}|\neg\text{spam}) P(\neg\text{spam})}$$

If  $P(\text{spam}|\text{free}) > P(\neg\text{spam}|\text{free})$ , we classify the email as **spam**.



# Naïve Bayes Classifier Preliminaries (II)

- What if we have more than one attribute?

$$P(y|X_1, X_2, \dots, X_p) = \frac{P(X_1, X_2, \dots, X_p|y) P(y)}{P(X_1, X_2, \dots, X_p)}$$

- Note that:

$$\begin{aligned} P(X_1, X_2, \dots, X_p|y) &= \frac{P(X_1, X_2, \dots, X_p, y)}{P(y)} \\ &= \frac{P(X_1, X_2, \dots, X_p, y)}{P(X_2, X_3, \dots, X_p, y)} \cdot \frac{P(X_2, X_3, \dots, X_p, y)}{P(X_3, X_4, \dots, X_p, y)} \cdot \dots \cdot \frac{P(X_p, y)}{P(y)} \\ &= P(X_1|X_2, \dots, X_p, y) P(X_2|X_3, \dots, X_p, y) \dots P(X_p|y) \end{aligned}$$

Let  $X_1$ ,  $X_2$ , and  $Y$  be random variables.  $X_1$  is **conditionally independent** of  $X_2$  given  $Y$  if

$$P(X_1|X_2, Y) = P(X_1|Y)$$





# Naïve Bayes Classifier Assumption

- **Naïve Bayes classifiers** assume that **all attributes are conditionally independent**. That is, if the class label  $y$  is known, then we consider the attributes to be **independent of each other**. That is:

$$P(X_1, X_2, \dots, X_p | y) = P(X_1 | y) P(X_2 | y) \dots P(X_p | y) = \prod_{i=1}^p P(X_i | y)$$

- Then, the probability of observing a class label  $y$  is given by:

$$P(y | X_1, X_2, \dots, X_p) = \frac{P(y) \prod_{i=1}^p P(X_i | y)}{P(X_1, X_2, \dots, X_p)}$$

- Since  $P(X_1, X_2, \dots, X_p)$  is constant for every class label, then:

$$P(y | X_1, X_2, \dots, X_p) \propto P(y) \prod_{i=1}^p P(X_i | y)$$

- Thus, it is sufficient to choose the class label that **maximizes**  $P(y) \prod_{i=1}^p P(X_i | y)$ .



# Naïve Bayes Classifier

## Advantages and Disadvantages


- What are some of the **advantages** of **Naïve Bayes**?
  - **Descriptive value.** Provides a probabilistic model of the relationship between the class label and the attributes, as well as probabilities that quantify the uncertainty in predictions.
  - Applicable to **categorical and numeric attributes**.
  - **Robust to irrelevant attributes.**
  - **Robust to noise.**
- What are some of the **disadvantages** of **Naïve Bayes**?
  - **Sensitive to redundant attributes.**
  - Assumes that **attributes are conditionally independent**. This assumption often does not hold in practice.
- We can represent other forms of conditional independence among attributes using **Bayesian Networks**, which are **probabilistic graphical models** that do not assume that attributes are conditionally independent.



# Support Vector Machines Preliminaries (I)

- A **hyperplane** is given by the equation  $\mathbf{w}^T \mathbf{x} + b = 0$  where  $\mathbf{x}$  are the attribute values and  $\mathbf{w}$  and  $b$  are the parameters of the **hyperplane**.
- An observation can belong to either side of the **hyperplane** depending on the sign of  $\mathbf{w}^T \mathbf{x} + b$ .
- A dataset is **linearly separable** if there exists a **hyperplane** that can perfectly separate observations from different classes.

if  $\mathbf{w}^T \mathbf{x}_i + b > 0$ ,  
then  $y_i = -$




if  $\mathbf{w}^T \mathbf{x}_i + b < 0$ ,  
then  $y_i = +$




# Classification Exercise 14.1




There can be infinitely many hyperplanes to separate the classes.  
Which of these hyperplanes would you choose? Why?



Option A



Option B



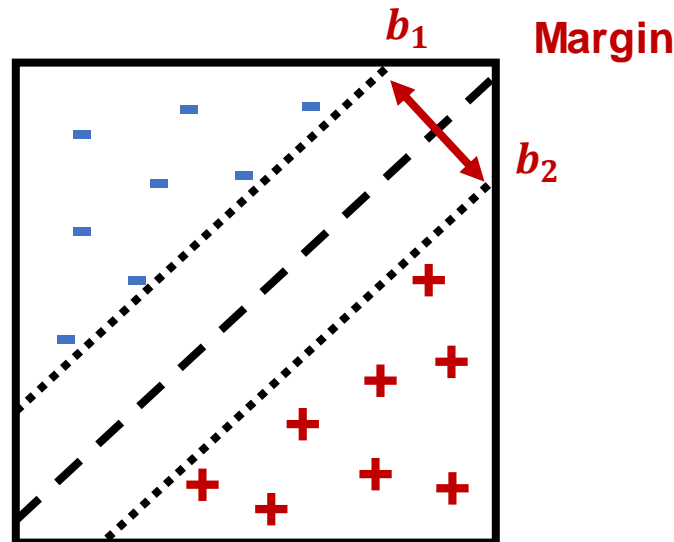
Option C



# Support Vector Machines Preliminaries (II)

- We would like to choose a **hyperplane** that is **robust to small perturbations in the data**. This can be achieved by choosing the **hyperplane** with the **maximum margin**.
- For every **hyperplane**  $B$ , there is a pair of **parallel hyperplanes**  $b_1$  and  $b_2$  that touch the **closest observations** of both classes.
- The **distance between  $b_1$  and  $b_2$**  is known as the **margin** of **hyperplane**  $B$ .

What happens if the margin is too small?





# Support Vector Machines


## Definition

- A **support vector machine (SVM)** is a classifier that finds the **separating hyperplane** with the **maximum margin** by solving the following **optimization problem**:

$$\begin{aligned} & \max_{w,b} M \\ & \text{subject to } y_i(w^T x_i + b) \geq M ||w|| \end{aligned}$$

where  $w$  and  $b$  are the parameters of the **hyperplane**,  $2M$  is the margin, and  $y_i \in \{-1, 1\}$ .

- SVMs** represent this **hyperplane** using only a subset of the observations in the training set that are **most difficult to classify** (known as the **support vectors**).
- If the number of **support vectors** is **too large**, the **SVM** is more likely to **overfit**.






# Support Vector Machines

## Nonlinear SVM


- The previous definition of **SVM** learns a **linear decision boundary** to separate the classes.
- We can learn **nonlinear decision boundaries** by transforming the data from the original attribute space  $\mathbf{x}$  into a new space  $\phi(\mathbf{x})$  where a **linear hyperplane** can be used to separate the classes.



**Not linearly  
separable**

$$\phi(x) = (x, x^2)$$

→



**Linearly  
separable**



# Support Vector Machines

## Kernel Trick

- Computing the **transformation**  $\phi(\mathbf{u})$  for every observation  $\mathbf{u}$  can be **computationally expensive**.
- However, in order to solve the **SVM** optimization problem, we only need the **inner products**  $\langle \phi(\mathbf{u}), \phi(\mathbf{v}) \rangle$ .
- Therefore, we can use a **kernel function**  $K(\mathbf{u}, \mathbf{v})$  to compute the **inner products**  $\langle \phi(\mathbf{u}), \phi(\mathbf{v}) \rangle$  without explicitly computing the **transformations**  $\phi(\mathbf{u})$  and  $\phi(\mathbf{v})$ .

$$K(\mathbf{u}, \mathbf{v}) = \langle \phi(\mathbf{u}), \phi(\mathbf{v}) \rangle = f(\mathbf{u}, \mathbf{v})$$

- This method is known as the **kernel trick**.
- Some commonly used **kernel functions** include:
  - **Polynomial kernel:**  $K(\mathbf{u}, \mathbf{v}) = (\mathbf{u}^T \mathbf{v} + 1)^p$
  - **Radial Basis Function (RBF) kernel:**  $K(\mathbf{u}, \mathbf{v}) = e^{-||\mathbf{u}-\mathbf{v}||^2/(2\sigma^2)}$





# Support Vector Machines

## Advantages and Disadvantages

- What are some of the **advantages** of **SVMs**?
  - Uses optimization algorithms to find the **optimal decision boundary**.
  - **Flexibility**. Can learn both linear and nonlinear decision boundaries.
  - Applicable to **numeric and categorical attributes** (by using dummy variables).
  - **Generally robust to irrelevant attributes** (by assigning them a weight of zero).
  - **Generally robust to redundant attributes** (by assigning them a similar weight).
- What are some of the **disadvantages** of **SVMs**?
  - Training the model can be **computationally expensive**.



# Classification References

- Joel Grus. *Data Science from Scratch* (2015).
- Cathy O'Neil and Rachel Schutt. *Doing Data Science* (2013).
- Pang-Ning Tan, Michael Steinbach, Anuj Karpatne, Vipin Kumar. *Introduction to Data Mining* (2019).